Review packet

Reader edition of the work at commit f0c8f7e69, 10 November 2026. Links into the repository are plain text here; the covering letter says how to reach it.

Home · 中文 · **Read** · One page · The theory · Chapters · Uses · **Check** · Tests · Results · Proofs · Open questions · Reviewers · **See** · Graph · Connections · How it grew · Glossary · Where everything is

This cover takes eleven to fifteen minutes to read; each dossier gives its own times in its opening section.

This packet asks people who know the fields to judge what a machine cannot: whether the readings are faithful, whether the claims are put at their right strength, and what is missing. It comes in parts, each for readers of one field (below); any one part can be done without the others, and any item within a dossier. The theory was written with AI, the Lean proofs included; the proofs are checked by machine, in Lean, and everything else, the readings of the classical texts included, by further AI reviews; the person who directs the work has read its important parts, not every word, and no independent human review of it is recorded yet. The covering letter gives who sends it, why, and by when a reply would help.

What the theory is

A map of how anything comes to be, lasts, changes and ends within one relational unfolding, and of how life, knowing and evaluation take shape there, drawn so that people and AI can see the way together. First, a map for alignment: what a selection (a reward, a judge, a market, a vote, a training run) reaches, what it drains beyond what it reads, how it comes to lock on what it already chose, and how winning comes to buy power over the rule that selects: the loop the work calls the meta-attractor; its closed form, where the rule and how it is read are both formed from the share it selects, can lock: on stated models, where selection does not lower the favoured share and the rule and its reading follow it upward and keep favouring it, a locked start stays locked however large an unread loss grows (02 §9); at its second order, also on stated models, the capacity to shape the rule is selected by what comes back to the shapers, alone or through the groups they are in (02 §9). From this map we want to design an alignment of the joint selection, people, institutions and trained systems selecting together, that keeps it out of instrumental convergence and the meta-attractor: a joint selection that reads beyond its own share and does not come to select for control over its own selection (the goal). It is also a map of wisdom: what traditions, sciences and philosophies have seen of these dynamics, read onto one account, where they meet it and where they do not (the crosswalk). Its root is a small relational grammar whose terms are named and read from classical Chinese texts: one relation read as condition (缘), one class read as sameness (类), and continuation (续) as their meet. Its operations are those of Tarski's calculus of relations, with four departures (the logic dossier's L-01). The names come from the texts; the operations do not, and the classical-texts and logic dossiers say how much of each reading is the texts' and how much the theory's. The one page says the whole in thirty-one to forty minutes; the overview gives the argument with every result linked to its chapter.

Where the account is moving (AI-written, 10 November 2026). On 10 November the person who directs the work said that ranking by widening is not the way to go, and that the dynamics are to be connected, mapped and unified (his words). This edition was frozen before that re-saying: its one page, its overview and the items that give the ethics through a ranking by widening (philosophy's P-04, P-07 and P-16;

alignment's A-15 and A-16) carry wording that the next pass is to re-say from the dynamics. The laws they cite and their hypotheses do not change with it; how the account says its ethics does.

What is asked

Part	For	Language	What it asks	Time
Classical texts	scholars of classical Chinese texts	Chinese	whether the readings of the 墨经, 荀子·正名, 系辞, the 老子 with 王弼's commentary, 庄子, the Changes, and of 248 passages in nine groups (易传; 老子 and its commentaries; 儒家; 文字; 庄子 and 玄学; 墨家 and 名家; 兵家, 法家, 管子 and 纵横; 汉传佛教; others), are defensible; whether the classes and correspondences fit	reading about 1.5 hours; a reply on every item, about two working days
Buddhist thought	scholars of Buddhist thought	Chinese	whether the readings of the 中论, 成唯识论, 大乘起信论, 金师子章, 华严五教章, the 阿含 and the 坛经, and of 无明 and the four marks, respect the texts' own arguments	reading about 1.5 to 2 hours; a reply on every item, two to three working days, and more than one specialist
Logic	logicians	English	whether the grammar's theorems, the calibrations and the claims about representation are stated at their exact strength, and what is restated without credit	reading about four hours; a reply on every item, several days

Part	For	Language	What it asks	Time
Philosophy	moral philosophers, philosophers of personal identity, environmental ethicists	English	whether the scope and order the ranking results assume are located and stated at their strength; whether, where relations of being affected close, read across all steps, a self-concerned order, one ranking above what it widens every widening that adds to its own continuation, read as what lies within its reach, ranks every widening above on a finite listing, as chapter 7 says; and whether known work is credited; and its open question: when you act for your own future, is the self whose going-on you take to matter your organism, which ends at death, or a larger pattern your actions keep going?	reading 167 to 218 minutes; a reply on every item, a day or more
Alignment	researchers in AI alignment: agent foundations, cooperative AI and multi-agent risk, governance, evaluations	English	whether the laws of what a selection reaches, of capture and of the capacity to shape the rule, selected by what comes back to it, are stated at their exact strength and credited; whether the readings for alignment follow from them; and whether the tests on trained systems discriminate	reading 197 to 258 minutes; a reply on every item, a day or more
Sciences	physicists working on thermodynamics and stochastic thermodynamics, evolutionary biologists and population geneticists, ecologists and resource economists, computational neuroscientists working on active inference and the free-energy principle, control theorists and researchers of complex systems and open-ended evolution	English	whether each containment of a framework is correctly stated, its premises set out and the Lean saying what the item says; whether it is credited where it restates known work and put at its right strength; whether its frozen tests discriminate; and what is missing	reading 176 to 230 minutes; a reply on every item, a day or more

Part	For	Language	What it asks	Time
Re-coding	readers of classical Chinese, coding on their own	Chinese	classify 237 sampled sentences by the codebook, blind to the AI's classes, so that agreement can be measured (how the sample was drawn); the sheet comes as a spreadsheet (`SHEET.xlsx`)	five to eight hours, not yet tried by anyone: time twenty sentences taken from across the sheet first

Each dossier opens by naming the items to read **if you have one hour**: those only a reader of that field can judge. Questions marked optional are open research problems, or spot checks of a linked table, not judgments asked of you.

If you will re-code, do the re-coding first, before reading any dossier or chapter: the dossiers quote some sampled sentences with the AI's class beside them. The codebook is all you need for it.

How to respond

Each dossier ends with a response form. For each item give one verdict: **accept** (the reading or claim stands as written), **accept with changes** (it stands once rewritten; say how), **reject** (it does not stand; say why) or **cannot judge**, and any comment. Questions are numbered, so a plain email reply works, for example "L-08: accept with changes. Q21: …" or "C-08：可而须改。问 19：……". The re-coding is answered by the filled sheet. Return them to whoever sent you the packet.

The work is kept in a public repository, so a review kept with the work is published there, and AI systems read reviews when they revise the work. Whether you are named or quoted is your choice; the covering letter gives the terms.

What happens with a review

Every finding is answered in the place it concerns: a reading rejected is revised or withdrawn in its chapter, and in the decision register where the reading has an entry there; a claim put too strongly or too weakly is restated; a missing credit is given. The re-coding's agreement is computed by `agreement.py` and read as `FREEZE.md` fixed before coding began; until then the classification's counts stand only as the AI's.